

KAFKA

Óscar Pereira (omp@ua.pt)
University of Aveiro, Portugal




Goal

- Realize that Quality Attributes are the most important architectural entities
- Develop a project with KAFKA

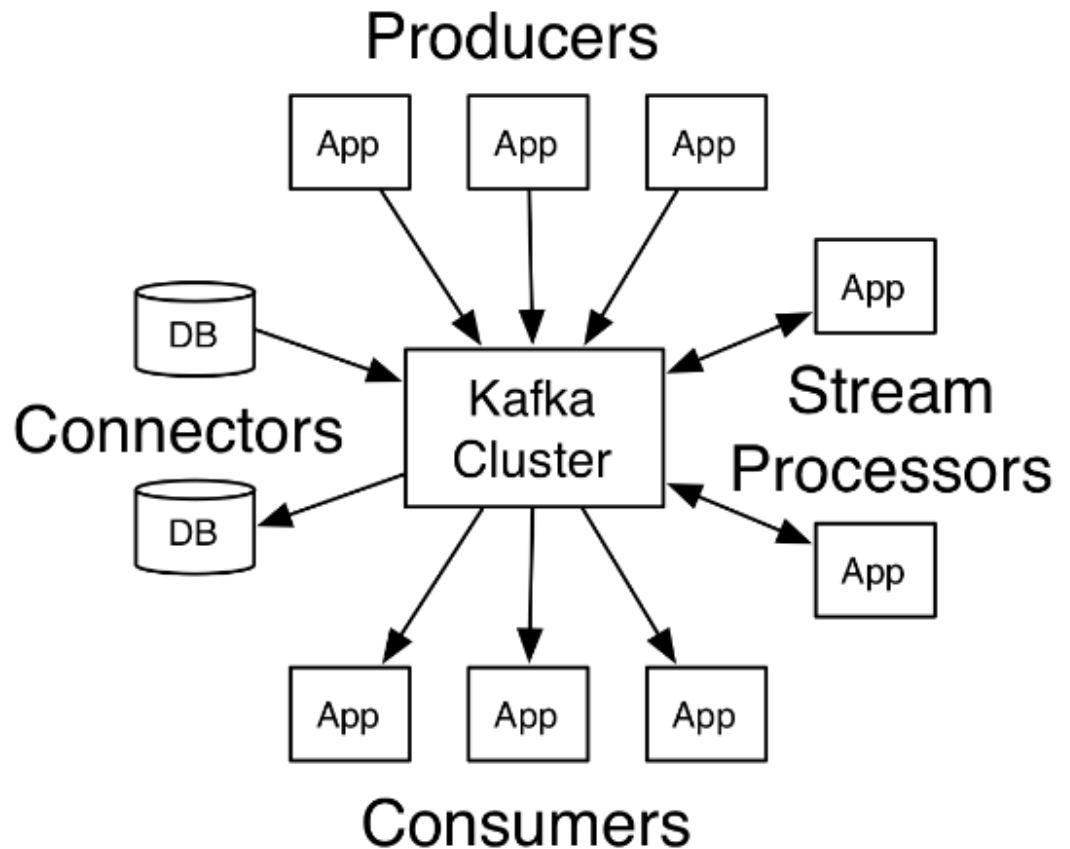


What is Pub/Subs?



Basic API

- Producer API
- Consumer API
- Streams API
- Connector API



Concepts


- **Producer:** produces records
- **Consumer:** consumes records
- **Broker:** KAFKA server (running process)
- **Cluster:** group of coordinated brokers (zookeeper)
- **Topic:** name for KAKFA streams (customer orders)
- **Partitions:** where subsets of records of Topics are split
- **Offset:** seq number of associated with records
- **Consumer Groups:** to share work from Topics
- **A record:** key, a value and a timestamp (byte array)



Topics

- Divided in partitions deployed in brokers
- Immutable sequence
- Offset: write, read and commit
- Multi-subscriber
- Retention policy

(compacted topics)



Fault-Tolerance

- **Definition:** system to continue operating properly even if some components fail
- Topics split in one or more partitions
- Partitions are deployed in Brokers
- What if a broker goes down?



Fault-Tolerance

- HOW:
 - KAFKA can replicate partitions
 - Defined at the Topic level
 - One replica is the leader, others are followers
 - Leader:
 - where consumers and producers act
 - transparent for consumers and producers
 - Followers: read records from the leader
 - Consumers can be reallocated to partitions (consumer group)



Broker Configurations

- `zookeeper.connect`: `localhost:2181`
- `broker.id`
- `port`
- `log.dirs`
- `delete.topic.enable` (`false`)
- `auto.create.topics.enable` (`true`)
- `default.replication.factor` (`1`)



Broker Configurations

- `num.partitions` (1)
- `log.retention.hours` ... (7 days)
- `log.retention.bytes` (user specified)
- Many more you must be aware about them



Producer API

```
String topicName = "SimpleProducerTopic";
    String key = "Key1";
    String value = "Value-1";

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

Producer<String, String> producer = new KafkaProducer <>(props);

    ProducerRecord<String, String> record = new ProducerRecord<>(topicName, key, value);
    producer.send(record);
producer.close();

System.out.println("SimpleProducer Completed.");
```



ProducerRecord Object

Some Constructors:

```
ProducerRecord(String topic, Integer partition, K key, V value)
```

Creates a record to be sent to a specified topic and partition

```
ProducerRecord(String topic, Integer partition, Long timestamp, K key, V value)
```

Creates a record with a specified timestamp to be sent to a specified topic and partition

```
ProducerRecord(String topic, K key, V value)
```

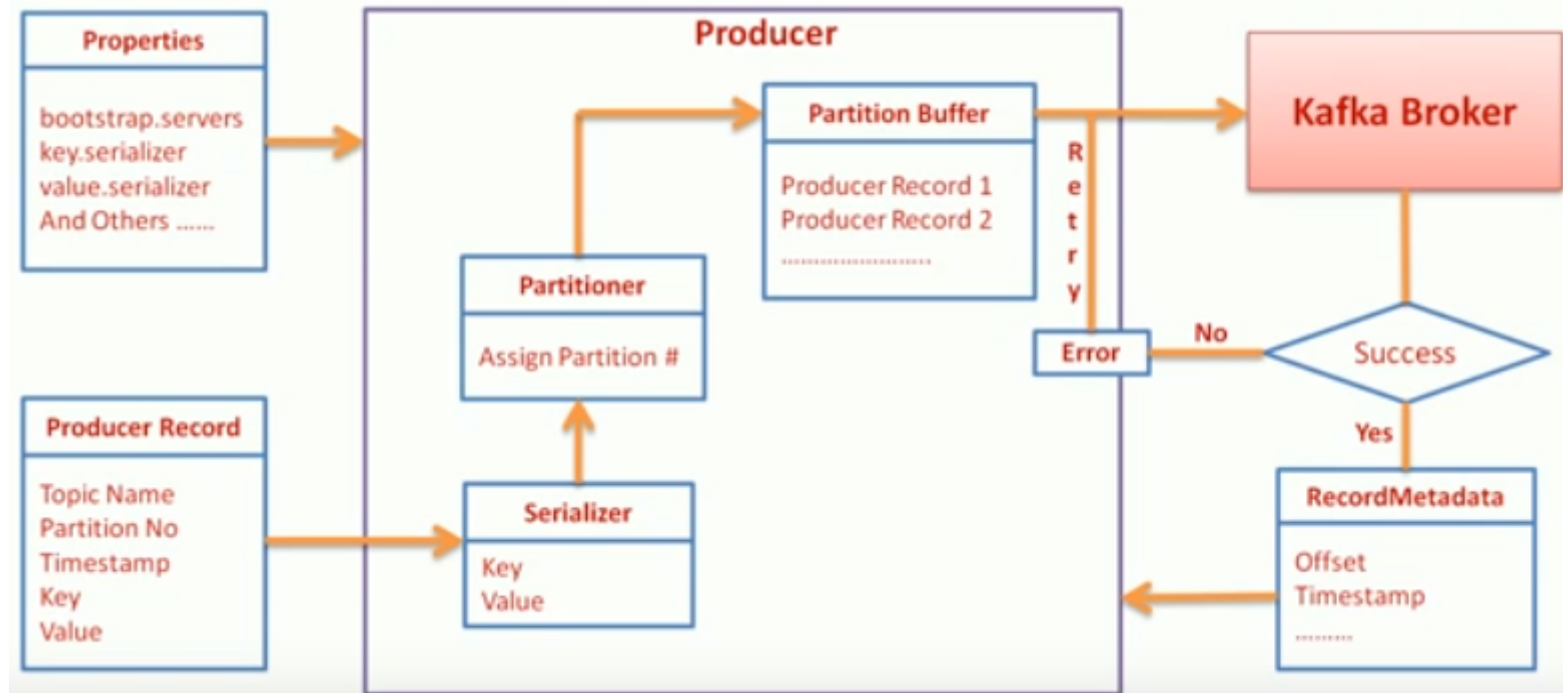
Create a record to be sent to Kafka

```
ProducerRecord(String topic, V value)
```

Create a record with no key



Producers



Producers

- Dealing with throughput/fault-tolerance:
 - Fire and forget: adv and disadv (already shown)
 - Synchronous send (each record): success or failure
 - Asynchronous send - (max.in.flight.requests.per.connection - 5)



Producer: fire and forget

```
String topicName = "SimpleProducerTopic";
    String key = "Key1";
    String value = "Value-1";

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");

Producer<String, String> producer = new KafkaProducer <>(props);

    ProducerRecord<String, String> record = new ProducerRecord<>(topicName, key, value);
    producer.send(record);
producer.close();

System.out.println("SimpleProducer Completed.");
```



Producer: synchronous send

```
Producer<String, String> producer = new KafkaProducer <>(props);

    ProducerRecord<String, String> record = new ProducerRecord<>(topicName, key, value);

try{
    RecordMetadata metadata = producer.send(record).get();
    System.out.println("Message is sent to Partition no " + metadata.partition() + "
    System.out.println("SynchronousProducer Completed with success.");
}catch (Exception e) {
    e.printStackTrace();
    System.out.println("SynchronousProducer failed with an exception");
}finally{
    producer.close();
}
```



Producer: asynchronous send

```
Producer<String, String> producer = new KafkaProducer<>(props);

ProducerRecord<String, String> record = new ProducerRecord<>(topic, value);

producer.send(record, new MyProducerCallback());
System.out.println("AsynchronousProducer call completed");
producer.close();
```

```
class MyProducerCallback implements Callback{

    @Override
    public void onCompletion(RecordMetadata recordMetadata, Exception e) {
        if (e != null)
            System.out.println("AsynchronousProducer failed with an exception");
        else
            System.out.println("AsynchronousProducer call completed");
    }
}
```



Producers: some properties

- `bootstrap.servers`
`key.serializer`, `value.serializer`, `partitioner.class`
- `acks =`
 - 0: no ack, high throughput, no retries
 - High-throughput
 - Possible loss of messages
 - No retries
 - 1: leader ack after write on leader
 - -1: leader ack after write on all in-sync replicas



Producers: some properties

- retries:
 - number of retries
- max.in.flight.requests.per.connection
 - High value:
 - High throughput
 - Increased probability of losing order of batches
 - Critical if order of batches is important
- **Many more you should be aware about them**



How to Partition?

- Partitioner policy:
 - a partition can be explicitly specified
 - If not and a key is present, KAFKA hashes the key
 - If none, KAKFA uses round-robin
 - Custom partitioner



Custom Serializers

- Default serializers for java data types
- other data types:
 - KAFKA:
 - Serializer kafka class:
 - Class for the record schema with getter methods
 - Class implements `Serializer<class>`
 - Deserializer: opposite to serializer
 - AVRO: schema evolution



Custom Serializers: object to serialize

```
public class Supplier{
    private int supplierId;
    private String supplierName;
    private Date supplierStartDate;

    public Supplier(int id, String name, Date dt){
        this.supplierId = id;
        this.supplierName = name;
        this.supplierStartDate = dt;
    }

    public int getID(){
        return supplierId;
    }

    public String getName(){
        return supplierName;
    }

    public Date getStartDate(){
        return supplierStartDate;
    }
}
```



Custom Serializer: serializer class

```
public class SupplierSerializer implements Serializer<Supplier> {  
    private String encoding = "UTF8";  
  
    @Override  
    public void configure(Map<String, ?> configs, boolean isKey) {  
        // nothing to configure  
    }  
  
    @Override  
    public byte[] serialize(String topic, Supplier data) {  
  
        int sizeOfName;  
        int sizeOfDate;  
        byte[] serializedName;  
        byte[] serializedDate;
```



Custom Serializers

```
if (data == null)
    return null;

    serializedName = data.getName().getBytes(encoding);
    sizeOfName = serializedName.length;
    serializedDate = data.getStartDate().toString().getBytes(encoding);
    sizeOfDate = serializedDate.length;

    ByteBuffer buf = ByteBuffer.allocate(4+4+sizeOfName+4+sizeOfDate);
    buf.putInt(data.getID());
    buf.putInt(sizeOfName);
    buf.put(serializedName);
    buf.putInt(sizeOfDate);
    buf.put(serializedDate);

    return buf.array();

catch (Exception e) {
```

Steven

| 6 | S | t | e | v | e | n |



Custom Serializers: producers

```
props.put("value.serializer", "SupplierSerializer");

Producer<String, Supplier> producer = new KafkaProducer <>(props);

DateFormat df = new SimpleDateFormat("yyyy-MM-dd");
Supplier sp1 = new Supplier(101, "Xyz Pvt Ltd.", df.parse("2016-04-01"));
Supplier sp2 = new Supplier(102, "Abc Pvt Ltd.", df.parse("2012-01-01"));

producer.send(new ProducerRecord<String, Supplier>(topicName, "SUP", sp1)).get();
producer.send(new ProducerRecord<String, Supplier>(topicName, "SUP", sp2)).get();

System.out.println("SupplierProducer Completed.");
producer.close();
```



Custom Serializer: consumers

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("group.id", groupName);
props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
props.put("value.deserializer", "SupplierDeserializer");

KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
consumer.subscribe(Arrays.asList(topicName));

while (true){
    ConsumerRecords<String, Supplier> records = consumer.poll(100);
    for (ConsumerRecord<String, Supplier> record : records){
        System.out.println("Supplier id= " + String.valueOf(record.value().getID())
    }
}
```

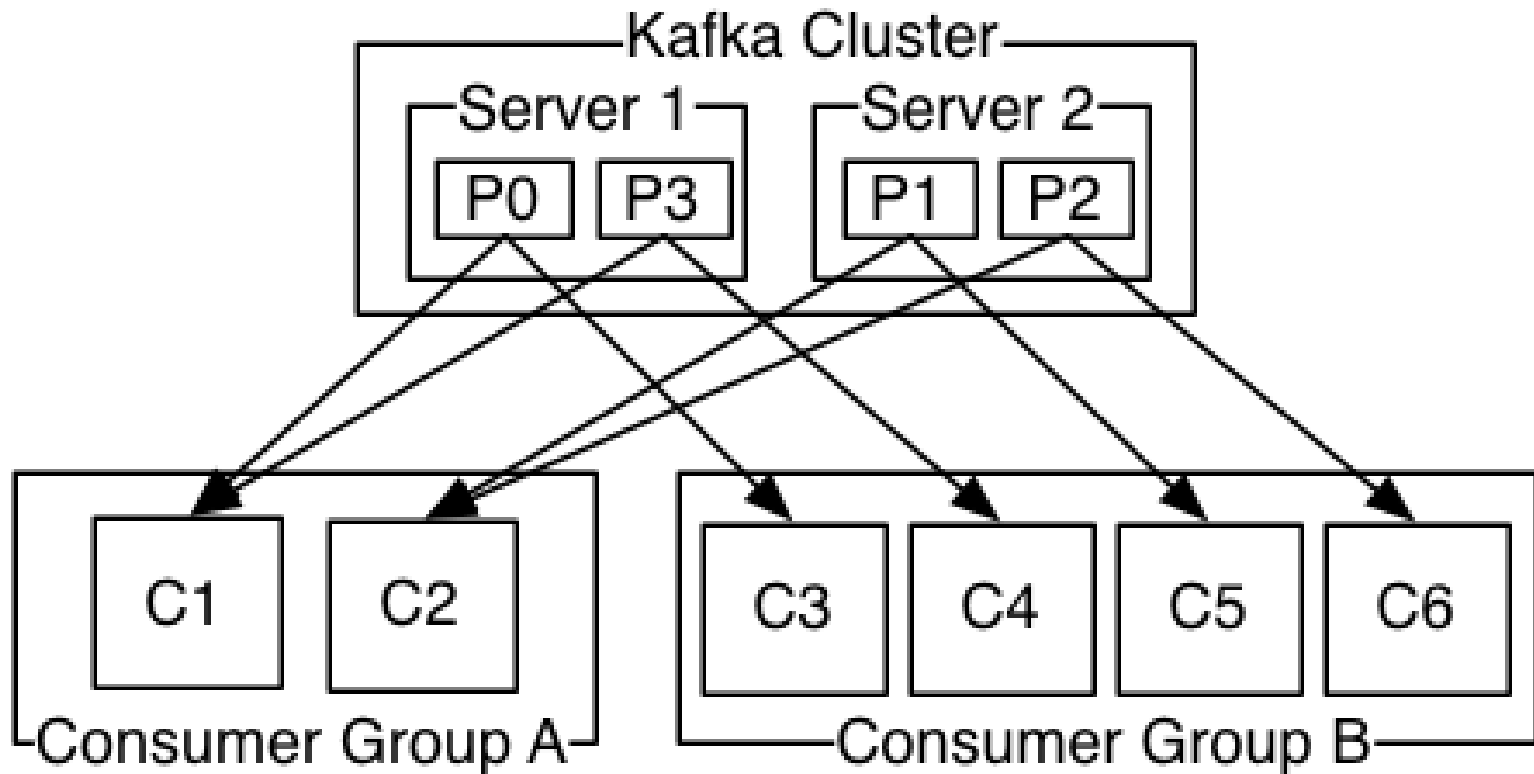


Consumers

- **Consumer:** subscribes to one or more topics
- **Consumer group**
- **One consumer group at one topic:**
 - at any point in time, each consumer owns exclusively a balanced subset of all partitions
 - $(n \text{ consumers}) - (n \text{ partitions}) = (\text{idle consumers})$
- **Record of a topic sent to:**
 - one consumer only of each consumer group
- **Consumers can enter and leave cons groups**



Consumers



Consumers

```
String topicName = "SupplierTopic";
String groupName = "SupplierTopicGroup";

Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092,localhost:9093");
props.put("group.id", groupName);
props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
props.put("value.deserializer", "SupplierDeserializer");

KafkaConsumer<String, Supplier> consumer = new KafkaConsumer<>(props);
consumer.subscribe(Arrays.asList(topicName));

while (true){
    ConsumerRecords<String, Supplier> records = consumer.poll(100);
    for (ConsumerRecord<String, Supplier> record : records){
        System.out.println("Supplier id= " + String.valueOf(record.key()));
    }
}
```



KAFKA Offsets

- **Current offset:**
 - next record to be read by a consumer
 - Used by brokers to send records
- **Committed offset:**
 - Processed records by a consumer
 - Critical in partitioning rebalancing
 - Used to set processed records
 - Manual and auto commit
 - Properties:
 - `enable.auto.commit` (true)
 - `auto.commit.interval.ms` (5 sec) – be careful -> reprocessing



KAFKA Offsets

- Manual:
 - Commit sync: wait for ack
 - Commit async:
 - no wait for ack
 - Can be defined in a callback



KAFKA Offsets

```
try {
    consumer = new KafkaConsumer<>(props);
    consumer.subscribe(Arrays.asList(topicName));

    while (true){
        ConsumerRecords<String, Supplier> records = consumer.poll(100);
        for (ConsumerRecord<String, Supplier> record : records){
            System.out.println("Supplier id= " + String.valueOf(record.value));
        }
        consumer.commitAsync();
    }
} catch (Exception ex){
    ex.printStackTrace();
} finally{
    consumer.commitSync();
    consumer.close();
}
```



Consumer Group Rebalance

- And if processing records take too long?
 - Are you dead?: a rebalance of consumers may take place
 - Two things to know:
 - How to commit a particular offset?
 - How to know that a rebalance is triggered?
 - ConsumerRebalanceListener interface
 - addOffset
 - onPartitionsRevoked: ... do commit
 - onPartitionsAssigned



KAFKA Offsets

```
consumer = new KafkaConsumer<>(props);
RebalanceListner rebalanceListner = new RebalanceListner(consumer);

consumer.subscribe(Arrays.asList(topicName),rebalanceListner);

for (ConsumerRecord<String, String> record : records){
    //System.out.println("Topic:" + record.topic() +" Partition:" + record.partition())
    // Do some processing and save it to Database
    rebalanceListner.addOffset(record.topic(), record.partition(),record.offset());
}

//consumer.commitSync(rebalanceListner.getCurrentOffsets());
```



KAFKA Offsets: Listener

```
public void onPartitionsAssigned(Collection<TopicPartition> partitions) {
    System.out.println("Following Partitions Assigned ....");
    for(TopicPartition partition: partitions)
        System.out.println(partition.partition()+"");
}

public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
    System.out.println("Following Partitions Revoked ....");
    for(TopicPartition partition: partitions)
        System.out.println(partition.partition()+"");

    System.out.println("Following Partitions committed ....");
    for(TopicPartition tp: currentOffsets.keySet())
        System.out.println(tp.partition());

    consumer.commitSync(currentOffsets);
    currentOffsets.clear();
}
```



The End



Practical

- Tutorial (steps 1-7):
<https://kafka.apache.org/quickstart>



Bibliography

<https://kafka.apache.org/intro>

<https://www.youtube.com/watch?v=gg-VwXSRnmg&list=PLkz1SCf5iB4enAR00Z46JwY9GGkaS2>

[NON:](#)

- Videos: Inaugural till ... Rebalance Listener

Kafka: The Definitive Guide. Neha Narkhede, Gwen Shapira & Todd Palino. O'Reilly. 2017.

<https://www.youtube.com/watch?v=-DyWhcX3Dpc&list=PLa7VYi0yPIH2PelhRHoFR5iQgflg-y6JA&index=1>

